

CS231n

Convolutional Neural Networks for Visual Recognition

Sungchan Yi

September 13, 2019

Contents

1	Introduction	2
2	Image Classification	2
3	Loss Functions and Optimization	5
4	Introduction to Neural Networks	8

1 Introduction

1.1 History of Computer Vision

- ...
- ImageNet: Image classification challenge
- Huge improvement on 2012 - Convolutional Neural Networks (**CNN**)

1.2 CS231n Overview

- Focus on **image classification**
- Object detection, action classification, image captioning ...
- **CNN** !

2 Image Classification

2.1 Image Classification

- Map: Input Image \rightarrow Predetermined Category
- Semantic Gap - computers only see the pixel values of the image
- Many problems that cause the pixel values to change
 - Viewpoints
 - Illumination
 - Deformation (poses)
 - Occlusion (part of the object)
 - Background Clutter (object looks similar to the background)
 - Interclass Variation (objects come in different sizes and shapes)
- Image classification algorithms must be able to handle all these problems

2.2 Attempts to Classify Images

- Finding corners or edges to determine the object
 - Brittle, doesn't work very well

- **Not scalable** to work with other objects

- **Data-Driven Approach**

1. Collect a dataset of images and labels
2. Use ML algorithms to train a classifier
3. Evaluate the classifier on new images

- Data-driven approach is much more general

2.3 Nearest Neighbor

1. Memorize all data and labels
2. Predict the label of the **most similar** training image

- How to compare images? **Distance Metric** !

- L_1 distance (Manhattan Distance)¹

$$d_1(I_1, I_2) = \sum_{\text{pixel } p} |I_1^p - I_2^p|$$

- With N examples, training takes $\mathcal{O}(1)$, prediction takes $\mathcal{O}(N)$.
- Generally, we want prediction to be fast, but slow training time is OK
- Only looking at a **single** neighbor may cause problems (outliers, noisy data, etc.)
- Motivation for **kNNs**

2.4 k -Nearest Neighbors

- Take the **majority vote** from k closest points
- L_2 distance (Euclidean Distance)

$$d_2(I_1, I_2) = \sqrt{\sum_{\text{pixel } p} (I_1^p - I_2^p)^2}$$

- L_1 distance depends on the coordinate system L_2 doesn't matter
- Generalizing the distance metric can lead to classifying other objects such as texts

¹Dumb way to compare, but does reasonable things *sometimes* ...

2.5 Hyperparameters

- **Hyperparameters** are choices about the algorithm that we set, rather than by learning the parameter from data
- In kNNs, the value of k and the distance metric are hyperparameters
- **Problem dependent.** Must go through trial and error to choose the best hyperparameter
- Setting Hyperparameters
 - Split the data into 3 sets. Training set, validation set, and test set
 - Train the algorithm with many different hyperparameters on the training set
 - Evaluate with validation set, choose the best hyperparameter from the results
 - Run it once on the test set, and this result goes on the paper
- Never used on image data ...
 - Slow on prediction
 - Distance metrics on pixels are not informative (Distance metric may not capture the difference between images)
 - Curse of dimensionality - data points increase exponentially as dimension increases, but there may not be enough data to cover the area densely (Need to densely cover the regions of the n -dimensional space, for the kNNs to work well)

2.6 Linear Classification

- Lego blocks - different components, as building blocks of neural networks
- **Parametric Approach**
 - Image x , weight parameters W for each pixel in the image
 - A function $f : (x, W) \rightarrow$ numbers giving class scores for each class
 - If input x is an $n \times 1$ vector and there were c classes, W must be $c \times n$ matrix
- Classifier summarizes our knowledge on the data and store it inside the parameter W
- At test time, we use the parameter W to predict
- How to come up with the right structure for f ?
- **Linear Classifier** : $f(x, W) = Wx (+ b)$
- There may be a bias term b to show preference for some class label

- (Idea) Each row of W will work as a template for matching the input image, and the dot product of each row and the input image vector will give the **similarity** of the input data to the class
- Problems with linear classifier
 - Learning single template for each class
 - If the class has variations on how the class might appear, the classifier averages out all the variations and tries to recognize the object by a single template
 - Using deeper models to remove this restriction will lead to better accuracy
- Linear classifiers draws hyperplanes on the n -dimensional space to classify images
- If hyperplanes cannot be drawn on the space, the linear classifier may struggle (parity problem, multi-modal data)

3 Loss Functions and Optimization

3.1 Motivation

- We saw that W can act as a template for each class
- How do we choose such W ?
- To choose the best W , we should be able to **quantify** the goodness of prediction across the training data

3.2 Loss Functions

- Dataset of examples: $\{(x_i, y_i)\}_{i=1}^N$, where x_i is an input data, and $y_i \in \mathbb{Z}$ is a correct label for the data
- Suppose we have a prediction function f , then the **loss over the dataset** is a sum of loss over examples

$$L = \frac{1}{N} \sum_i L_i(f(x_i, W), y_i)$$

- Now we want to choose a W that **minimizes the loss function**
- **Multiclass SVM loss**
 - Scores vector $s = f(x_i, W)$
 - SVM loss for each data (Hinge loss)

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

- As the score for the true category (s_{y_i}) increases, the loss goes down linearly, until it gets above a safety margin because now the example is correctly classified
- Simply put - We are happy¹ if *the score for the correct label is much higher than all the other scores* by some margin²
- Minimum loss is 0, maximum loss is ∞
- Quadratic hinge loss - May be used to put more loss on scores that are totally off (Depends on how we want to weigh off different mistakes that the classifier makes)
- Suppose we found a W that makes $L = 0$. But this W is not unique³, so how do we choose such W ?
- We have only written down loss **in terms of the data**. We only told the classifier to find the W that fits the data
- But in practice, we only care about the performance on the test data

3.3 Regularization

- We add an additional **regularization** term to the loss function, which **encourages** the model to somehow pick a simpler W
- We use a regularization penalty(loss) $R(W)$, then

$$L = \frac{1}{N} \sum_i L_i(f(x_i, W), y_i) + \lambda R(W)$$

and the λ is a hyperparameter that trades off between data loss and regularization loss

- Common regularization functions

- L_2 regularization

$$R(W) = \sum_i \sum_j W_{ij}^2$$

- L_1 regularization

$$R(W) = \sum_i \sum_j |W_{ij}|$$

- Anything that you do to the model that *penalizes the complexity of the model*
- L_1 regularization thinks less non-zero entries are less complex, L_2 thinks spreading numbers all across entries of W is less complex

¹We are happy when the loss is small

²The constant 1 in the equation can actually be generalized

³ W will also give 0 loss

3.4 Softmax Loss

- Multinomial Logistic Regression
- Multiclass SVM - no interpretation for each score
- Consider the scores as *unnormalized log probabilities of the classes*
- $s = f(x_i, W)$

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}} \quad (\text{softmax function})$$

- Now we have a probability distribution, and we want this to match the distribution that put all its weight on the correct label
- Loss is the $-\log$ of the probability of the true class⁴ (**Cross-Entropy**)

$$L_i = -\log \left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \right)$$

- Minimum loss is 0, maximum loss is ∞

3.5 Optimization

- Minimize the loss function!
- Strategy #1 : Random Search (...)
- Strategy #2 : **Follow the slope** - Use the geometry of the space to find *which direction from here will decrease the loss function*
- **Gradients** ! - Directional derivatives to find out which direction we should take a step, to minimize the loss function
- Use calculus to compute analytic gradients and use them in code⁵

3.6 Gradient Descent

- The direction of the negative gradient is the fastest decrease
- Key Idea: Starting from the initial guess x_0 , iteratively compute x_n by the following

$$x_{n+1} = x_n - \gamma \cdot \nabla F(x_n)$$

and hope to converge to some x

⁴You can view this as the KL divergence between the target and computed distribution, maximum likelihood estimate

⁵Analytic gradients are fast, exact and error-prone, while numerical one is slow and approximate

- Constant γ is the **step size**, sometimes called the **learning rate**, which is an important hyper-parameter
- This is **slow** ... Computing gradients for each iteration is too costly
- **Stochastic Gradient Descent** uses a **minibatch** (subset of training data) to estimate the true gradient

3.7 Image Features

- Feeding raw pixel values don't work well
- Compute various **feature representations** of the image, concatenate them and feed it to the classifier
- What is the best feature transform? ⁶
- Example: Color Histogram, Histogram of Oriented Gradients (Local orientation of edges)

4 Introduction to Neural Networks

4.1 Computational Graphs

- Graph to represent any function, where each node is a function (some operation)
- Advantage - **Backpropagation**: Recursively using the chain rule in order to compute the gradient with respect to every variable in the computational graph

4.2 Backpropagation

- Suppose we have a computational graph for loss function L , and we try to calculate the gradient at some node corresponding to the operation f
- Suppose f takes x, y as inputs and outputs z , i.e. $z = f(x, y)$
- We are given the gradient for z , namely $\frac{\partial L}{\partial z}$. With this we want to calculate $\frac{\partial L}{\partial x}$ and $\frac{\partial L}{\partial y}$
- Use the **chain rule** to get

$$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial z} \cdot \frac{\partial z}{\partial x} \quad \text{and} \quad \frac{\partial L}{\partial y} = \frac{\partial L}{\partial z} \cdot \frac{\partial z}{\partial y}$$

⁶For instance, conversion from polar to Cartesian coordinates

- Usually operation f is simple, so we can find the local gradients, $\frac{\partial z}{\partial x}$ and $\frac{\partial z}{\partial y}$ ^{1 2}
- Now $\frac{\partial L}{\partial x}$ and $\frac{\partial L}{\partial y}$ will be used **recursively** for the nodes that produced x, y

- **Sigmoid function**

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Useful fact is that

$$\frac{d\sigma(x)}{dx} = (1 - \sigma(x)) \cdot \sigma(x)$$

- Patterns in backward flow
 - Addition gate *distributes* gradients
 - Max gate *routes* gradients
 - Multiplication gate *switches* gradients
- Note that multiple gradients coming in from upstream are added together
- When x, y, z become vectors, the gradients will be **Jacobian matrices**
- But computational graphs will be very large, so it is impractical to write down each gradient formula by hand for all parameters

4.3 Neural Networks

- We talked about linear score function $f = Wx$
- 2-layer neural network $f = W_2 \cdot \max(0, W_1x) \dots$
- Stack multiple layers with non-linear function in between³
- **Activation functions**

¹With this method, we won't have to derive an expression for the gradient for cases where the loss function is very complicated. We can just work on the level where gradients are easy to compute

²The expression for $\partial L / \partial z$ may be complex, but in practice, this value will be a single number received from upstream, which we multiply to the local gradient. The expressions for local gradients will be all we need, which will be very simple

³Composition of linear functions will simplify to a linear function